

05_pandas

November 9, 2022

1 pandas

pandas is an open source library that brings high-performance, easy-to-use data structures and analysis tools to the Python programming language. Its core features are its 1-dimensional **Series** and 2-dimensional **DataFrame** objects. The former is like a NumPy array (**pandas** is built on top of NumPy) with an explicit index, and the latter is a bit like an Excel worksheet.

The [10 minutes to pandas](#) and [getting started](#) guides are excellent places to learn more, but read on for my own tour of the library which draws comparisons between **pandas** and familiar spreadsheet programs (e.g., [Microsoft Excel](#), [Apple Numbers](#), [LibreOffice Calc](#), [Google Sheets](#)).

1.1 Importing pandas

The community agreed custom for importing **pandas** is:

```
import pandas as pd
```

This gives us access to all **pandas** functionality, which means it is a bit like clicking on the desktop icon that launches a spreadsheet application. Let's import **pandas**, and while we are at it, let's also import **numpy** and **matplotlib** (and set some configuration options).

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('bmh')
```

We'll begin with a detailed look at the **Series** and **DataFrame** objects.

1.2 Series

pd.Series is much like a 1-dimensional NumPy array, but with explicit axis labels, the ability to hold any data type, and various feature and method enhancements to simplify working with complex data. We can think of a **Series** as the rough equivalent of a column in a spreadsheet application.

The basic way of making a **Series** is as follows:

```
s = pd.Series(data, index=index)
```

Note that the argument for `data` can take various forms, such as a Python dictionary, a `numpy` array, a list, or even a single value. If provided, `index` must be a list of axis labels with the same length as `data`.

Here's a simple `Series` of random numbers.

```
[2]: s = pd.Series(np.random.random(100))
      print(s)
```

```
0      0.110035
1      0.398505
2      0.783878
3      0.776024
4      0.343782
...
95     0.153074
96     0.580834
97     0.695021
98     0.302373
99     0.391739
Length: 100, dtype: float64
```

We didn't specify an `index`, so a default integer-based index (starting at 0) was applied, which we can see to the left of the actual values. Below the index and values we can also see the length and data type.

Now, let's make a `Series` containing the letters of the alphabet and assign also an `index` and a `name`.

```
[3]: letters = list('abcdefghijklmnopqrstuvwxyz')
      s = pd.Series(letters, index=range(1, 27), name='lowercase')
      print(s)
```

```
1      a
2      b
3      c
4      d
5      e
6      f
7      g
8      h
9      i
10     j
11     k
12     l
13     m
14     n
15     o
16     p
17     q
```

```

18    r
19    s
20    t
21    u
22    v
23    w
24    x
25    y
26    z
Name: lowercase, dtype: object

```

A few things to note here. First, the `index` doesn't default to zero because we specified our own labels. Second, the `name` we provided is shown below the values (this is analogous to a column header in spreadsheet applications). Third, the data type is `object`, denoting that the `Series` contains `str` data.

1.3 DataFrame

A `DataFrame` is a 2-dimensional data structure consisting of rows and columns, which makes it similar to a worksheet in Microsoft Excel. Each column in a `DataFrame` is a `Series`, which by now we know to be a feature-enhanced, explicitly indexed, NumPy array that can contain any type of data. A `DataFrame`, then, is a collection of `Series` that share a common `Index`.

An easy way to create a basic `DataFrame` is to pass a Python dict to the `pd.DataFrame(...)` constructor. In this scenario, the keys will serve as names for the columns, and the values, which must be convertible to a series-like data structure, will be the data.

```

[4]: # It's the familiar shopping example again...
df = pd.DataFrame(
    {
        'item': ['apples', 'bananas', 'bread', 'yoghurt', 'pasta', 'coffee',
        ↪ 'butter', 'carrots'],
        'unit_price': np.array([2.20, 1.85, 1.10, 1.20, 1.45, 2.80, 2.40, .89]),
        ↪ dtype='float32'),
        'quantity': np.array([2, 1, 2, 4, 2, 1, 1, 1]),
        'date_of_purchase': pd.Timestamp('20221208'),
        'shop': pd.Categorical(['Aldi'] * 8),
        'paid_cash': True
    }
)
df

```

```

[4]:      item  unit_price  quantity  date_of_purchase  shop  paid_cash
0  apples         2.20         2    2022-12-08  Aldi        True
1  bananas         1.85         1    2022-12-08  Aldi        True
2   bread         1.10         2    2022-12-08  Aldi        True
3  yoghurt         1.20         4    2022-12-08  Aldi        True
4   pasta         1.45         2    2022-12-08  Aldi        True
5   coffee         2.80         1    2022-12-08  Aldi        True

```

6	butter	2.40	1	2022-12-08	Aldi	True
7	carrots	0.89	1	2022-12-08	Aldi	True

As with NumPy arrays, the `DataFrame` object stores some basic information as attributes. Below, we get information about shape, number of dimensions, size (number of elements), and the data types for each column of `df`.

```
[5]: print('Shape: ', df.shape)
      print('Number of dimensions: ', df.ndim)
      print('Number of elements: ', df.size)
      print('Data types:\n')
      print(df.dtypes)
```

```
Shape: (8, 6)
Number of dimensions: 2
Number of elements: 48
Data types:
```

```
item          object
unit_price    float32
quantity      int64
date_of_purchase  datetime64[ns]
shop          category
paid_cash     bool
dtype: object
```

We can also see how much memory (in bytes) is being used by each column.

```
[6]: df.memory_usage()
```

```
[6]: Index          128
      item           64
      unit_price     32
      quantity       64
      date_of_purchase  64
      shop          124
      paid_cash       8
      dtype: int64
```

The actual values stored in a `DataFrame` are accessible via the `.values` attribute. Inspection of these values shows that `pandas` is built right on top of NumPy.

```
[7]: print("The values of the data frame are: ")
      print(df.values)
      print("The type of the values is ", type(df.values))
```

The values of the data frame are:

```
[['apples' 2.2000000047683716 2 Timestamp('2022-12-08 00:00:00') 'Aldi'
  True]
```

```
['bananas' 1.850000023841858 1 Timestamp('2022-12-08 00:00:00') 'Aldi'
 True]
['bread' 1.100000023841858 2 Timestamp('2022-12-08 00:00:00') 'Aldi'
 True]
['yoghurt' 1.2000000476837158 4 Timestamp('2022-12-08 00:00:00') 'Aldi'
 True]
['pasta' 1.4500000476837158 2 Timestamp('2022-12-08 00:00:00') 'Aldi'
 True]
['coffee' 2.799999952316284 1 Timestamp('2022-12-08 00:00:00') 'Aldi'
 True]
['butter' 2.4000000953674316 1 Timestamp('2022-12-08 00:00:00') 'Aldi'
 True]
['carrots' 0.8899999856948853 1 Timestamp('2022-12-08 00:00:00') 'Aldi'
 True]]
```

The type of the values is `<class 'numpy.ndarray'>`

1.4 Viewing, selecting and indexing

If you are new to `pandas`, it may sometimes feel like the data are hidden away and out of reach, especially in comparison to spreadsheet environments where data are generally always on the screen and can be selected and browsed by clicking and scrolling the mouse. Thankfully, `pandas` has some handy tools for viewing a `DataFrame`.

`DataFrame.head()` can be used to view the top of a frame, and `DataFrame.tail()` to view the bottom. The default number of rows to display is 5, but we can change this by specifying a number.

```
[8]: # Show the first three rows of df
df.head(3)
```

```
[8]:      item  unit_price  quantity  date_of_purchase  shop  paid_cash
0  apples         2.20         2      2022-12-08  Aldi      True
1  bananas         1.85         1      2022-12-08  Aldi      True
2   bread         1.10         2      2022-12-08  Aldi      True
```

```
[9]: # Show the last three rows of df
df.tail(3)
```

```
[9]:      item  unit_price  quantity  date_of_purchase  shop  paid_cash
5  coffee         2.80         1      2022-12-08  Aldi      True
6  butter         2.40         1      2022-12-08  Aldi      True
7  carrots         0.89         1      2022-12-08  Aldi      True
```

The row and column indices themselves can also be accessed and operated upon (note that we didn't specify a row index at the time of creation, so we got a default `RangeIndex`).

```
[10]: # View the row index of df
df.index
```

```
[10]: RangeIndex(start=0, stop=8, step=1)
```

```
[11]: # View the column index of df
df.columns
```

```
[11]: Index(['item', 'unit_price', 'quantity', 'date_of_purchase', 'shop',
          'paid_cash'],
          dtype='object')
```

Let's talk about selecting columns. To select a single column of a `DataFrame`, put its name in square brackets. This is a bit like clicking on the column header in a spreadsheet to highlight the entire column. The result of this operation is a `Series`.

```
[12]: df['unit_price']
```

```
[12]: 0    2.20
      1    1.85
      2    1.10
      3    1.20
      4    1.45
      5    2.80
      6    2.40
      7    0.89
      Name: unit_price, dtype: float32
```

As long as the column name doesn't contain whitespace or special characters, single columns may also be accessed in equivalent fashion via `.` notation. I often find this more convenient, as it's easier to type.

```
[13]: df.unit_price
```

```
[13]: 0    2.20
      1    1.85
      2    1.10
      3    1.20
      4    1.45
      5    2.80
      6    2.40
      7    0.89
      Name: unit_price, dtype: float32
```

Selecting multiple columns at the same time requires a list of column names, and returns a `DataFrame` (this is like doing `Ctrl+click` or `Ctrl+drag` to select multiple columns in a spreadsheet).

```
[14]: df[['item', 'quantity', 'unit_price']]
```

```
[14]:   item  quantity  unit_price
0  apples         2         2.20
1  bananas         1         1.85
2   bread         2         1.10
```

3	yoghurt	4	1.20
4	pasta	2	1.45
5	coffee	1	2.80
6	butter	1	2.40
7	carrots	1	0.89

[] can be used to slice by row

```
[15]: df[2:5]
```

```
[15]:      item  unit_price  quantity  date_of_purchase  shop  paid_cash
2    bread      1.10         2      2022-12-08    Aldi      True
3  yoghurt      1.20         4      2022-12-08    Aldi      True
4    pasta      1.45         2      2022-12-08    Aldi      True
```

pandas also supports **label-based indexing**, which is extremely useful for selecting individual values or cross-sections of a DataFrame.

Label-based indexing is performed using the `.loc[]` and `.at[]` access methods. Note these are square, and not round, brackets.

```
[16]: # Select rows 4:7 and the specified columns
df.loc[4:7, ['item', 'unit_price', 'quantity']]
```

```
[16]:      item  unit_price  quantity
4    pasta      1.45         2
5    coffee      2.80         1
6    butter      2.40         1
7  carrots      0.89         1
```

```
[17]: # Select the value in row=3, column='item'
df.at[3, 'item']
```

```
[17]: 'yoghurt'
```

Positional indexing can be achieved with the `.iloc[]` and `.iat[]` access methods. The difference here is that we **must** provide integers (starting at 0). The behavior is very similar to array slicing in Python and NumPy.

```
[18]: # Select the 6th row
df.iloc[5]
```

```
[18]: item                                coffee
unit_price                                2.8
quantity                                  1
date_of_purchase    2022-12-08 00:00:00
shop                                Aldi
paid_cash                                True
Name: 5, dtype: object
```

```
[19]: # Select the value in the second row of the first column
df.iat[1, 0]
```

```
[19]: 'bananas'
```

[Boolean indexing](#) is another common way to select a subset of data from a frame. Suppose we cared only about rows where the `unit_price` was over £2.00.

```
[20]: df[df['unit_price'] > 2.]
```

```
[20]:
```

	item	unit_price	quantity	date_of_purchase	shop	paid_cash
0	apples	2.2	2	2022-12-08	Aldi	True
5	coffee	2.8	1	2022-12-08	Aldi	True
6	butter	2.4	1	2022-12-08	Aldi	True

This works, because the expression inside the square brackets evaluates as a `Series` of `bool`, and the indexing operation ensures that we only get the rows where the expression comes up `True`.

```
[21]: df['unit_price'] > 2.
```

```
[21]: 0    True
      1   False
      2   False
      3   False
      4   False
      5    True
      6    True
      7   False
      Name: unit_price, dtype: bool
```

We can get quite fancy by chaining multiple expressions together. For example, if we only wanted rows where the `unit_price` is greater than £1.00 and less than £2.00.

```
[22]: df[((df['unit_price'] > 1.) & (df['unit_price'] < 2.))]
```

```
[22]:
```

	item	unit_price	quantity	date_of_purchase	shop	paid_cash
1	bananas	1.85	1	2022-12-08	Aldi	True
2	bread	1.10	2	2022-12-08	Aldi	True
3	yoghurt	1.20	4	2022-12-08	Aldi	True
4	pasta	1.45	2	2022-12-08	Aldi	True

1.5 World population dataset

The best way to learn `pandas` is to use it with real data, so let's explore some more advanced features with an actual dataset.

In the `data` folder of the course materials, there's a file called `world_population.csv` which I sourced from an excellent data science website called [Kaggle](#). The file contains historical population

data for every country/territory in the world, along with various other parameters. Here's what's in the file:

Variable	Definition
Rank	Rank by Population
CCA3	3 Digit Country/Territories Code
Country	Name of the Country/Territories
Capital	Name of the Capital
Continent	Name of the Continent
2022 Population	Population of the Country/Territories in the year 2022
2020 Population	Population of the Country/Territories in the year 2020
2015 Population	Population of the Country/Territories in the year 2015
2010 Population	Population of the Country/Territories in the year 2010
2000 Population	Population of the Country/Territories in the year 2000
1990 Population	Population of the Country/Territories in the year 1990
1980 Population	Population of the Country/Territories in the year 1980
1970 Population	Population of the Country/Territories in the year 1970
Area (km ²)	Area size of the Country/Territories in square kilometer
Density (per km ²)	Population Density per square kilometer
Growth Rate	Population Growth Rate by Country/Territories
World Population Percentage	The population percentage by each Country/Territories

Let's start by loading and inspecting the data.

Because its a CSV file, `pd.read_csv(...)` is right tool for the job.

```
[23]: df = pd.read_csv('../data/world_population.csv')
df
```

```
[23]:
```

	Rank	CCA3	Country	Capital	Continent	\
0	36	AFG	Afghanistan	Kabul	Asia	
1	138	ALB	Albania	Tirana	Europe	
2	34	DZA	Algeria	Algiers	Africa	
3	213	ASM	American Samoa	Pago Pago	Oceania	
4	203	AND	Andorra	Andorra la Vella	Europe	
..	
229	226	WLF	Wallis and Futuna	Mata-Utu	Oceania	
230	172	ESH	Western Sahara	El Aaiún	Africa	
231	46	YEM	Yemen	Sanaa	Asia	
232	63	ZMB	Zambia	Lusaka	Africa	
233	74	ZWE	Zimbabwe	Harare	Africa	

	2022 Population	2020 Population	2015 Population	2010 Population	\
0	41128771	38972230	33753499	28189672	
1	2842321	2866849	2882481	2913399	
2	44903225	43451666	39543154	35856344	
3	44273	46189	51368	54849	

4	79824	77700	71746	71519
..
229	11572	11655	12182	13142
230	575986	556048	491824	413296
231	33696614	32284046	28516545	24743946
232	20017675	18927715	16248230	13792086
233	16320537	15669666	14154937	12839771

	2000 Population	1990 Population	1980 Population	1970 Population	\
0	19542982	10694796	12486631	10752971	
1	3182021	3295066	2941651	2324731	
2	30774621	25518074	18739378	13795915	
3	58230	47818	32886	27075	
4	66097	53569	35611	19860	
..	
229	14723	13454	11315	9377	
230	270375	178529	116775	76371	
231	18628700	13375121	9204938	6843607	
232	9891136	7686401	5720438	4281671	
233	11834676	10113893	7049926	5202918	

	Area (km ²)	Density (per km ²)	Growth Rate	World Population Percentage
0	652230	63.0587	1.0257	0.52
1	28748	98.8702	0.9957	0.04
2	2381741	18.8531	1.0164	0.56
3	199	222.4774	0.9831	0.00
4	468	170.5641	1.0100	0.00
..
229	142	81.4930	0.9953	0.00
230	266000	2.1654	1.0184	0.01
231	527968	63.8232	1.0217	0.42
232	752612	26.5976	1.0280	0.25
233	390757	41.7665	1.0204	0.20

[234 rows x 17 columns]

We now have a `DataFrame` with 17 columns and 234 rows, which appear to be sorted by `Country` in ascending alphabetical order. All in all, these are clean and well organized data, but if we sorted the rows by `Rank`, we could get a quick and easy insight into the most and least populous countries. We can do this using `.sort_values(...)`.

```
[24]: # Sort the rows by Rank in descending order
df.sort_values('Rank', ascending=False)
```

```
[24]:
```

	Rank	CCA3	Country	Capital	Continent	\
41	1	CHN	China	Beijing	Asia	
92	2	IND	India	New Delhi	Asia	

221	3	USA	United States	Washington, D.C.	North America
93	4	IDN	Indonesia	Jakarta	Asia
156	5	PAK	Pakistan	Islamabad	Asia
..
137	230	MSR	Montserrat	Brades	North America
64	231	FLK	Falkland Islands	Stanley	South America
150	232	NIU	Niue	Alofi	Oceania
209	233	TKL	Tokelau	Nukunonu	Oceania
226	234	VAT	Vatican City	Vatican City	Europe

	2022 Population	2020 Population	2015 Population	2010 Population	\
41	1425887337	1424929781	1393715448	1348191368	
92	1417173173	1396387127	1322866505	1240613620	
221	338289857	335942003	324607776	311182845	
93	275501339	271857970	259091970	244016173	
156	235824862	227196741	210969298	194454498	
..	
137	4390	4500	5059	4938	
64	3780	3747	3408	3187	
150	1934	1942	1847	1812	
209	1871	1827	1454	1367	
226	510	520	564	596	

	2000 Population	1990 Population	1980 Population	1970 Population	\
41	1264099069	1153704252	982372466	822534450	
92	1059633675	870452165	696828385	557501301	
221	282398554	248083732	223140018	200328340	
93	214072421	182159874	148177096	115228394	
156	154369924	115414069	80624057	59290872	
..	
137	5138	10805	11452	11402	
64	3080	2332	2240	2274	
150	2074	2533	3637	5185	
209	1666	1669	1647	1714	
226	651	700	733	752	

	Area (km ²)	Density (per km ²)	Growth Rate	World Population Percentage
41	9706961	146.8933	1.0000	17.88
92	3287590	431.0675	1.0068	17.77
221	9372610	36.0935	1.0038	4.24
93	1904569	144.6529	1.0064	3.45
156	881912	267.4018	1.0191	2.96
..
137	102	43.0392	0.9939	0.00
64	12173	0.3105	1.0043	0.00
150	260	7.4385	0.9985	0.00
209	12	155.9167	1.0119	0.00

226	1	510.0000	0.9980	0.00
-----	---	----------	--------	------

[234 rows x 17 columns]

This is great, but everyone knows that China is the most populated country in the world. Surely there is a more nuanced story to be told. Let's start by calculating the most recent estimate of the global population by summing all of the values in the 2022 Population column.

```
[25]: print('The global population in 2022 is: ')
      df['2022 Population'].sum()
```

The global population in 2022 is:

```
[25]: 7973413042
```

That is believable. But now I'm curious how that is distributed by Continent.

A `.groupby` operation is perfect for this job.

[Read more about `.groupby\(\)`](#)


```
[26]: # Group by Continent and sum the values
      df.groupby('Continent')['2022 Population'].sum()
```

```
[26]: Continent
      Africa      1426730932
      Asia       4721383274
      Europe      743147538
      North America 600296136
      Oceania      45038554
      South America 436816608
      Name: 2022 Population, dtype: int64
```

Unsurprisingly, Asia is the most highly populated continent. A bar chart may help us to appreciate this more fully. We can easily make a bar chart using the `.plot()` method for `Series` and `DataFrames`.


[Read more about `.plot\(\)`](#)

```
[27]: # Repeat the above operation, but this time chain a plotting method to the end
      (
      df.groupby('Continent')['2022 Population']
      .sum()
      .plot(kind='bar', rot=45, xlabel='', ylabel='Population')
      );
```



So Asia is the most populated continent, but surely that's just because it's the biggest. How does the population relate to the total size of the continent? The **Density (per km²)** column looks like it was calculated by dividing the population of a country by its total area, so repeating the operation with this variable will help to develop the story.

```
[28]: # As above, but with the 'Density (per km2)' column instead
(
    df.groupby('Continent')['Density (per km2)']
      .sum()
      .plot(kind='bar', rot=45, xlabel='', ylabel='Population density (per km2)')
);
```



This paints a different picture. Most striking is the difference for Europe. Though it is the third most populated continent, Europe comes second for population density, and by a long way in comparison to the others. This is especially true in comparison with South America, whose population is very sparse indeed.

1.6 The complete picture - population growth

To tell the overall story of these population data, I came up with my own two-figure solution, which I feel does a pretty good job. I sort the data by **Rank**, calculate a column-wise differential on the **<date> Population** columns, and then plot the result for each country in a stacked horizontal bar chart, with bars to the left show population decline, and bars to right show population growth for that period. To avoid ambiguity, I put the total population values to the side of each bar. On the right, these values indicate the population for that country in 2022, and on the left, they indicate the total population decline since 1970.

```
[33]: # Get the columns with the population data
pop_cols = df.columns[df.columns.str.endswith('Population')]

# Sort and select data
data = (df
```

```

.set_index('Country') # Set 'Country' as the index
.sort_values('Rank', ascending=True)[pop_cols] # Sort by 'Rank' and keep only
↳ the columns with population data
.iloc[0:117] # Choose 117 MOST populated countries with location-based
↳ indexing
.loc[:, lambda df_: reversed(df_.columns)] # Reverse the order of the columns
↳ in a fancy way
)

# Pull out the base population at 1970
base_pop_1970 = data['1970 Population']

# Calculate a column-wise differential. This means the data will
# now reflect change from the previous timepoint, rather than the
# total population at that time.
data = data.diff(axis=1)

# Put the base population back in
data['1970 Population'] = base_pop_1970

# Make a figure and axis that's big enough to show the data
fig, ax = plt.subplots(figsize=(12, 24))

# Plot a horizontal bar chart with pandas plotting method.
# Note the reversal of the data with[::-1], which is done
# to make sure the longest bars are at the top
data[::-1].plot(kind='barh', stacked=True, ax=ax)

# Add the numbers as text
for i, (country, row) in enumerate(data[::-1].iterrows()):
    pop = row.sum()
    text_pos = row[row>0].sum()
    ax.text(text_pos+1e7, i, '{:,}'.format(pop), va='center', fontsize=8)
    text_pos_neg = row[row < 0].abs().sum()
    if text_pos_neg > 0:
        ax.text(-(text_pos_neg+1e7), i, '{:,}'.format(text_pos_neg),
↳ va='center', ha='right', fontsize=8)

# Add a vertical black line at zero as a visual aid
ax.axvline(0, 0, 1, c='k', lw=.5)

# Format the x axis
def billions_formatter(x, pos):
    return f'{x / 1000_000_000}'
ax.xaxis.set_major_formatter(plt.FuncFormatter(billions_formatter))

# Tweak axis

```

```


ax.set(
    ylabel='',
    title='Population change in the 117 most populated countries/territories_
↪(1970-2022)',
    xlabel='Population in billions',
    xlim=(-.25e9, 1.75e9)
);

# Put legend in lower right of figure
ax.legend(loc='lower right')

# Save the figure with a tight bounding box and high resolution (300_
↪dots-per-inch)
fig.savefig('../images/most_populated_countries_2022.png', bbox_inches='tight',_
↪dpi=300)

```


Population change in the 117 most populated countries/territories (1970-2022)



Rinse and repeat for the other half of the data. By using a second figure for the less-populated half of the dataset, we reset the scale on the x-axis and dramatically improve the explanatory power of the visualisations.

```
[34]: # Get the columns with the population data
pop_cols = df.columns[df.columns.str.endswith('Population')]

# Sort and select data
data = (df
        .set_index('Country') # Set 'Country' as the index
        .sort_values('Rank', ascending=True)[pop_cols] # Sort by 'Rank' and keep only
        ↪ the columns with population data
        .iloc[-117:] # Choose the 117 LEAST populated countries with location-based
        ↪ indexing
        .loc[:, lambda df_: reversed(df_.columns)] # Reverse the order of the columns
        ↪ in a fancy way
    )

# Pull out the base population at 1970
base_pop_1970 = data['1970 Population']

# Calculate a column-wise differential. This means the data will
# now reflect change from the previous timepoint, rather than the
# total population at that time.
data = data.diff(axis=1)

# Put the base population back in
data['1970 Population'] = base_pop_1970

# Make a figure and axis that's big enough to show the data
fig, ax = plt.subplots(figsize=(12, 24))

# Plot a horizontal bar chart with pandas plotting method.
# Note the reversal of the data with[::-1], which is done
# to make sure the longest bars are at the top
data[::-1].plot(kind='barh', stacked=True, ax=ax)

# Add the numbers as text
for i, (country, row) in enumerate(data[::-1].iterrows()):
    pop = row.sum()
    text_pos = row[row > 0].sum()
    ax.text(text_pos+1e5, i, '{:,}'.format(pop), va='center', fontsize=8)
    text_pos_neg = row[row < 0].abs().sum()
    if text_pos_neg > 0:
```

```

        ax.text(-(text_pos_neg+1e5), i, '{:,}'.format(text_pos_neg),
↪va='center', ha='right', fontsize=8)

# Add a vertical line at zero as a visual aid
ax.axvline(0, 0, 1, c='k', lw=.5)

# Format the x axis
def millions_formatter(x, pos):
    return f'{x / 1_000_000}'
ax.xaxis.set_major_formatter(plt.FuncFormatter(millions_formatter))


# Tweak axis
ax.set(
    ylabel='',
    title='Population change in the 117 least populated countries/territories_
↪(1970-2022)',
    xlabel='Population in millions',
    xlim=(-3e6, 8e6)
)

# Put legend in lower right of figure
ax.legend(loc='lower right')

# Save the figure with a tight bounding box and high resolution (300_
↪dots-per-inch)
fig.savefig('../images/least_populated_countries_2022.png',
↪bbox_inches='tight', dpi=300)

```

Population change in the 117 least populated countries/territories (1970-2022)



1.7 The big picture - population density

After plotting the above, I realised I could do the same using only the Density (per km²) column. This was a little bit easier as it didn't require stacking the bars.

```
[35]: # Sort and select data
data = (df
        .set_index('Country') # Set 'Country' as the index
        .sort_values('Density (per km2)', ascending=True)
        .iloc[-117:] # Choose the 117 MOST densely populated countries with
        ↪ location-based indexing
    )

# Make a figure and axis that's big enough to show the data
fig, ax = plt.subplots(figsize=(12, 24))

# Plot a horizontal bar chart with pandas plotting method.
data.plot(kind='barh', y='Density (per km2)', ax=ax, legend=False)


# Add the numbers in text
for i, (country, row) in enumerate(data.iterrows()):
    pop = row['Density (per km2)']
    ax.text(pop+100, i, '{:,}'.format(round(pop, 2)), va='center', fontsize=8)

# Add a vertical line at zero as a visual aid
ax.axvline(0, 0, 1, c='k', lw=.5)

# Tweak axis
ax.set(
    ylabel='',
    title='Population density in the 117 most populated countries/territories',
    xlabel='Population (per km2)',
    xlim=(0, 25000)
)

# Save the figure with a tight bounding box and high resolution (300
    ↪ dots-per-inch)
fig.savefig('../images/most_dense_countries_2022.png', bbox_inches='tight',
    ↪ dpi=300)
```

Population density in the 117 most populated countries/territories



What's it like to be one of 23,172 people in a square kilometer of land? Go to Macau and you'll find out. Also, I didn't realise Gibraltar was so dense.

Now for the other half...

```
[36]: # Sort and select data
data = (df
    .set_index('Country') # Set 'Country' as the index
    .sort_values('Density (per km²)', ascending=True)
    .iloc[0:117] # Choose the 117 LEAST densely populated countries with
    ↪ location-based indexing
)

# Make a figure and axis that's big enough to show the data
fig, ax = plt.subplots(figsize=(12, 24))

# Plot a horizontal bar chart with pandas plotting method.
data.plot(kind='barh', y='Density (per km²)', ax=ax, legend=False)


# Add the numbers in text
for i, (country, row) in enumerate(data.iterrows()):
    pop = row['Density (per km²)']
    ax.text(pop+1, i, '{:,}'.format(round(pop, 2)), va='center', fontsize=8)

# Add a vertical line at zero as a visual aid
ax.axvline(0, 0, 1, c='k', lw=.5)

# Tweak axis
ax.set(
    ylabel='',
    title='Population density in the 117 least populated countries/territories',
    xlabel='Population (per km²)',
    xlim=(0, 100)
)

# Save the figure with a tight bounding box and high resolution (300
    ↪ dots-per-inch)
fig.savefig('../images/least_dense_countries_2022.png', bbox_inches='tight',
    ↪ dpi=300)
```


Population density in the 117 least populated countries/territories



Iceland... A beautiful country, where a kilometer of land is shared by only 3.62 people.

That's it for now. These examples are complicated, but if you study them carefully, change bits, and come up with your own variations, you will learn a lot about Python and **pandas** in the process! If you are feeling adventures, why not try and replicate these plots for the **Area (km²)** column? In this situation, countries like Russia and Canada would be right at the top, and at the bottom would be The Vatican City and small island nations like the Falkland and Faroe Islands.