

# 06\_more\_pandas

November 16, 2022

## 1 More about pandas

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('bmh')
```

We've looked at `Series`, `DataFrame`, how to view, select and index data, and how to make plots with `pandas`. Now we will look at time series data, which `pandas` is simply brilliant for working with.

'Time series data' refers to any form of data that is represented or indexed with ordered timestamps. Examples include stock prices, electrical activity in the brain, and temperature throughout the day. As you might have guessed, this involves dates, times, and differences between them.

### 1.1 Timestamp and Timedelta

In `pandas`, dates and times are represented with `pd.Timestamp`, which is a replacement for Python's native `datetime.datetime` object. The two are interchangeable in many respects, but a key difference is that `pd.Timestamp` uses NumPy's `datetime64` and `timedelta64` data types and incorporates a much wider range of functionality and powerful features for creating and manipulating time series.

Let's create timestamps to mark two memorable points in history: the end of the First and Second World Wars.

```
[33]: end_of_ww1 = pd.Timestamp('1918-11-11 11:00:00', ) # Armistice Day
end_of_ww2 = pd.Timestamp('1945-05-08 23:01:00') # VE Day
```

`Timestamp` data have many useful methods. An especially useful method is `.strftime()`, which creates a `str` representation of a timestamp using various formatting codes. Here are some examples.

```
[34]: print("The First World War ended", end_of_ww1.strftime('%d %B %Y'), "at",
      ↪end_of_ww1.strftime('%I:%M %p'))
print("The Second World War ended", end_of_ww2.strftime('%d %B %Y'), "at",
      ↪end_of_ww2.strftime('%I:%M %p'))
```

The First World War ended 11 November 1918 at 11:00 AM  
The Second World War ended 08 May 1945 at 11:01 PM

Here, the % symbols and subsequent characters tell the `strftime()` method how to represent the timestamp. You may wish to experiment with [other formatting codes](#) to format the timestamps in different ways. Copy the above code into the cell below and try to format the dates using abbreviated month names and 24-hour time.

```
[ ]:
```

If you are a history buff, you may already know the day of the week when each of The Great Wars ended. If not, you can find out using the `pd.Timestamp.day_name()` method.

```
[35]: print('The First World War ended on a', end_of_ww1.day_name())  
      print('The Second World War ended on a', end_of_ww2.day_name())
```

The First World War ended on a Monday  
The Second World War ended on a Tuesday

If we have two timestamps, we can calculate the difference between them, which is known as a `Timedelta`. Let's do this to find out how much time elapsed between the end of the First and Second World Wars.

```
[37]: tdelta = end_of_ww2 - end_of_ww1 # Subtract the end of WW1 from the end of WW2  
      print('Time between the end of the First and Second World Wars:')  
      print(f'{tdelta.days} days and {tdelta.seconds} seconds')
```

Time between the end of the First and Second World Wars:  
9675 days and 43260 seconds

A `Timedelta` behaves like a number in many respects and can be operated on with most mathematical operators. But there are restrictions. For example, one can divide or multiply a `Timedelta` by any number, but addition and subtraction can only take place with other `Timestamp` or `Timedelta` objects. If you think about it, this makes perfect sense.

```
[38]: print('Time between end of WW1 and WW2: ', tdelta)  
      print('Multiplied by 3: ', tdelta * 3)  
      print('Divided by 2: ', tdelta / 2)  
      print('Added to itself: ', tdelta + tdelta)  
      print('Minus twice itself: ', tdelta - (tdelta * 2))
```

Time between end of WW1 and WW2: 9675 days 12:01:00  
Multiplied by 3: 29026 days 12:03:00  
Divided by 2: 4837 days 18:00:30  
Added to itself: 19351 days 00:02:00  
Minus twice itself: -9676 days +11:59:00

If you try to perform an operation that is not supported, such as raising a `Timedelta` to some power or multiplying two `Timedelta`'s together, you will get an error.

```
[43]: tdelta ** 2 # Can not raise a Timedelta to a power...
```

```

-----
TypeError                                Traceback (most recent call last)
/var/folders/c9/7yddv1n2ss863cgfngj0wpm0000gp/T/ipykernel_7891/977053077.py in
    ↪<module>
----> 1 tdelta ** 2 # Can not raise a Timedelta to a power...

TypeError: unsupported operand type(s) for ** or pow(): 'Timedelta' and 'int'

```

In case you are wondering, leap years and other time-related idiosyncrasies are dealt with implicitly.

```

[48]: leap_year = pd.Timestamp('27 Feb 2020') # 2020 was a leap year
      print((leap_year + pd.Timedelta('2D')).strftime('%d %B %Y')) # Add 2 days

```

29 February 2020

```

[49]: not_a_leap_year = pd.Timestamp('27 Feb 2021') # 2021 wasn't a leap year
      print((not_a_leap_year + pd.Timedelta('2D')).strftime('%d %B %Y')) # Add 2 days

```

01 March 2021

## 1.2 Working with dates and times

Previously we have encountered functions such as `range()` and `np.linspace()`, which produce sequences of evenly-spaced numeric data. In `pandas`, we can do the same with dates using the `pd.date_range(...)` function. Let's use it to create a timestamp for every day of this year.

```

[50]: ts = pd.date_range(start='2022-01-01', end='2022-12-31', freq='D',
    ↪inclusive='both')
      ts

```

```

[50]: DatetimeIndex(['2022-01-01', '2022-01-02', '2022-01-03', '2022-01-04',
    '2022-01-05', '2022-01-06', '2022-01-07', '2022-01-08',
    '2022-01-09', '2022-01-10',
    ...,
    '2022-12-22', '2022-12-23', '2022-12-24', '2022-12-25',
    '2022-12-26', '2022-12-27', '2022-12-28', '2022-12-29',
    '2022-12-30', '2022-12-31'],
    dtype='datetime64[ns]', length=365, freq='D')

```

The result is a `DatetimeIndex`, which may be used as the index for a `Series` or `DataFrame`. Let's work some magic on these timestamps to get a list of all future dates for this year in long-format.

```

[51]: # Get all future dates and format them nicely
      ts[ts > pd.Timestamp.now()].strftime('%A %d %B %Y').tolist()

```

```

[51]: ['Thursday 17 November 2022',
    'Friday 18 November 2022',

```

```
'Saturday 19 November 2022',  
'Sunday 20 November 2022',  
'Monday 21 November 2022',  
'Tuesday 22 November 2022',  
'Wednesday 23 November 2022',  
'Thursday 24 November 2022',  
'Friday 25 November 2022',  
'Saturday 26 November 2022',  
'Sunday 27 November 2022',  
'Monday 28 November 2022',  
'Tuesday 29 November 2022',  
'Wednesday 30 November 2022',  
'Thursday 01 December 2022',  
'Friday 02 December 2022',  
'Saturday 03 December 2022',  
'Sunday 04 December 2022',  
'Monday 05 December 2022',  
'Tuesday 06 December 2022',  
'Wednesday 07 December 2022',  
'Thursday 08 December 2022',  
'Friday 09 December 2022',  
'Saturday 10 December 2022',  
'Sunday 11 December 2022',  
'Monday 12 December 2022',  
'Tuesday 13 December 2022',  
'Wednesday 14 December 2022',  
'Thursday 15 December 2022',  
'Friday 16 December 2022',  
'Saturday 17 December 2022',  
'Sunday 18 December 2022',  
'Monday 19 December 2022',  
'Tuesday 20 December 2022',  
'Wednesday 21 December 2022',  
'Thursday 22 December 2022',  
'Friday 23 December 2022',  
'Saturday 24 December 2022',  
'Sunday 25 December 2022',  
'Monday 26 December 2022',  
'Tuesday 27 December 2022',  
'Wednesday 28 December 2022',  
'Thursday 29 December 2022',  
'Friday 30 December 2022',  
'Saturday 31 December 2022']
```

As you can see, a `DatetimeIndex` is a powerful and flexible tool for representing dates and times. When it really starts to shine is when it is used as the index for time series data, so let's invent some.

Imagine there is a shop that opened at the turn of the millennium which sells apples, bananas and cakes. At the time of opening, these items were sold at the respective prices of £0.45, £0.28, and £0.62 per item, but there's been some inflation over time. Each day, a random(ish) number of each item is sold. Cakes are most popular, then bananas, then apples.

```
[53]: # Initial cost in year 2000
cost_of_apple = 0.45
cost_of_banana = 0.28
cost_of_cake = 0.62

# Create a range of dates from 2000 to 2022 with 1-day frequency.
# This means we will have a row for every day.
dates = pd.date_range('2000', '2022', freq='1D')
number_of_days = len(dates)

# Create the DataFrame
df = pd.DataFrame(
    {
        'date': dates,
        'apples_sold': np.random.randint(0, 200, number_of_days),
        'bananas_sold': np.random.randint(0, 350, number_of_days),
        'cakes_sold': np.random.randint(0, 700, number_of_days),
        'cost_of_apple': cost_of_apple,
        'cost_of_banana': cost_of_banana,
        'cost_of_cake': cost_of_cake
    },
)

# Adjust costs for inflation
cost_cols = ['cost_of_apple', 'cost_of_banana', 'cost_of_cake']
inflation = np.exp(np.linspace(0, 1, number_of_days))
df[cost_cols] = df[cost_cols].mul(inflation, axis=0)

# New columns for profit
df['apples_profit'] = df['apples_sold'] * df['cost_of_apple']
df['bananas_profit'] = df['bananas_sold'] * df['cost_of_banana']
df['cakes_profit'] = df['cakes_sold'] * df['cost_of_cake']
df['total_profit'] = df.apples_profit + df.bananas_profit + df.cakes_profit
df = df.round(2)
df
```

```
[53]:
```

	date	apples_sold	bananas_sold	cakes_sold	cost_of_apple	\
0	2000-01-01	175	174	534	0.45	
1	2000-01-02	61	316	394	0.45	
2	2000-01-03	164	320	487	0.45	
3	2000-01-04	55	291	143	0.45	
4	2000-01-05	8	261	630	0.45	

```

...
8032 2021-12-28      153      199      467      1.22
8033 2021-12-29      167       67      234      1.22
8034 2021-12-30       85      146      480      1.22
8035 2021-12-31       8       93      338      1.22
8036 2022-01-01       53     103       27      1.22

```

```

      cost_of_banana  cost_of_cake  apples_profit  bananas_profit \
0          0.28        0.62        78.75        48.72
1          0.28        0.62        27.45        88.49
2          0.28        0.62        73.82        89.62
3          0.28        0.62        24.76        81.51
4          0.28        0.62         3.60        73.12
...
8032          0.76        1.68       187.06       151.39
8033          0.76        1.68       204.20        50.98
8034          0.76        1.68       103.95       111.10
8035          0.76        1.69         9.78        70.78
8036          0.76        1.69        64.83        78.40

```

```

      cakes_profit  total_profit
0          331.08        458.55
1          244.31        360.25
2          302.02        465.46
3           88.69        194.96
4          390.79        467.51
...
8032          786.66       1125.11
8033          394.22        649.40
8034          808.76       1023.80
8035          569.57        650.13
8036           45.50        188.73

```

[8037 rows x 11 columns]

Now, if we look at the data types of the `DataFrame`, we can see that the `date` column has type `datetime64[ns]`.

```
[54]: df.dtypes
```

```

[54]: date          datetime64[ns]
      apples_sold      int64
      bananas_sold      int64
      cakes_sold       int64
      cost_of_apple     float64
      cost_of_banana     float64
      cost_of_cake       float64

```

```

apples_profit      float64
bananas_profit     float64
cakes_profit       float64
total_profit       float64
dtype: object

```

When working with the world population and titanic datasets, we used the `str` accessor method to interface with additional string methods (e.g., `df['name'].str.contains('Smith')`). For datetime-like data, there is a `dt` accessor method, which opens the door to datetime functionality. Below, we use the `dt` accessor method to create a separate column containing the day of the week, the year, and the fiscal year.

```

[55]: df['Day'] = df['date'].dt.day_name() # Get day of week
      df['Year'] = df['date'].dt.year # Get current year
      df['Fiscal Year'] = df['date'].dt.to_period('Q-APR').dt.qyear # Get fiscal year
      df.tail()

```

```

[55]:      date  apples_sold  bananas_sold  cakes_sold  cost_of_apple \
8032  2021-12-28          153           199          467           1.22
8033  2021-12-29          167            67          234           1.22
8034  2021-12-30           85          146          480           1.22
8035  2021-12-31           8           93          338           1.22
8036  2022-01-01           53          103           27           1.22

```

```

      cost_of_banana  cost_of_cake  apples_profit  bananas_profit \
8032             0.76           1.68         187.06          151.39
8033             0.76           1.68         204.20           50.98
8034             0.76           1.68         103.95          111.10
8035             0.76           1.69           9.78           70.78
8036             0.76           1.69          64.83           78.40

```

```

      cakes_profit  total_profit      Day  Year  Fiscal Year
8032         786.66       1125.11  Tuesday  2021       2022
8033         394.22        649.40 Wednesday  2021       2022
8034         808.76       1023.80  Thursday  2021       2022
8035         569.57        650.13   Friday  2021       2022
8036          45.50        188.73  Saturday  2022       2022

```

We now have three new and useful columns in our `DataFrame`. At the moment, the original `date` column is just a column of data, but what we really need is for it to be our index.

```

[59]: # Set the date column as the index
      df.index = df.date

```

Now we can index the `DataFrame` using any valid date/time string, of which there are many. For example, if we wanted to look at the information for a particular day, say, 23 April 2012, we could use any of the following:

- '2012-04-23'

- 'April 23, 2012'
- '23 Apr 2012'
- '23 April 2012'
- '2012/04/23'

```
[75]: df.loc['2012/04/23'] # St Georges day, 2012
```

```
[75]: date                2012-04-23 00:00:00
apples_sold              19
bananas_sold            322
cakes_sold              123
cost_of_apple            0.79
cost_of_banana           0.49
cost_of_cake             1.08
apples_profit            14.96
bananas_profit           157.76
cakes_profit             133.44
total_profit             306.16
Day                      Monday
Year                     2012
Fiscal Year              2012
Name: 2012-04-23 00:00:00, dtype: object
```

Try experimenting with different indexing strings!

```
[ ]:
```

If you want the data for an entire year, just pass in the year.

```
[80]: df.loc['2012'] # Get all data from 2012
```

```
[80]:
```

	date	apples_sold	bananas_sold	cakes_sold	cost_of_apple	\
date						
2012-01-01	2012-01-01	197	35	31	0.78	
2012-01-02	2012-01-02	4	338	212	0.78	
2012-01-03	2012-01-03	116	299	143	0.78	
2012-01-04	2012-01-04	155	309	566	0.78	
2012-01-05	2012-01-05	49	89	662	0.78	
...	...	...	...	...	...	
2012-12-27	2012-12-27	175	268	538	0.81	
2012-12-28	2012-12-28	150	101	349	0.81	
2012-12-29	2012-12-29	174	88	347	0.81	
2012-12-30	2012-12-30	133	206	599	0.81	
2012-12-31	2012-12-31	17	164	200	0.81	

	cost_of_banana	cost_of_cake	apples_profit	bananas_profit	\
date					
2012-01-01	0.48	1.07	152.95	16.91	



2012-01-02	0.48	1.07	3.11	163.31
2012-01-03	0.48	1.07	90.08	144.48
2012-01-04	0.48	1.07	120.39	149.33
2012-01-05	0.48	1.07	38.06	43.02
...	...	...	...	...
2012-12-27	0.51	1.12	142.11	135.42
2012-12-28	0.51	1.12	121.83	51.04
2012-12-29	0.51	1.12	141.34	44.48
2012-12-30	0.51	1.12	108.05	104.13
2012-12-31	0.51	1.12	13.81	82.91

	cakes_profit	total_profit	Day	Year	Fiscal Year
date					
2012-01-01	33.16	203.02	Sunday	2012	2012
2012-01-02	226.81	393.22	Monday	2012	2012
2012-01-03	153.01	387.57	Tuesday	2012	2012
2012-01-04	605.68	875.40	Wednesday	2012	2012
2012-01-05	708.50	789.58	Thursday	2012	2012
...	...	...	...	...	...
2012-12-27	601.95	879.48	Thursday	2012	2013
2012-12-28	390.53	563.40	Friday	2012	2013
2012-12-29	388.34	574.15	Saturday	2012	2013
2012-12-30	670.45	882.62	Sunday	2012	2013
2012-12-31	223.88	320.60	Monday	2012	2013

[366 rows x 14 columns]

Looks like 2012 was a leap year, because there are 366 days. Let's check this by getting the data for February of that year. To do this, we just need to add the month.

```
[88]: df.loc['2012-02'] # February 2012
```

```
[88]:
```

	date	apples_sold	bananas_sold	cakes_sold	cost_of_apple \
date					
2012-02-01	2012-02-01	85	41	271	0.78
2012-02-02	2012-02-02	178	281	337	0.78
2012-02-03	2012-02-03	36	134	47	0.78
2012-02-04	2012-02-04	191	81	223	0.78
2012-02-05	2012-02-05	51	176	314	0.78
2012-02-06	2012-02-06	69	265	190	0.78
2012-02-07	2012-02-07	129	154	20	0.78
2012-02-08	2012-02-08	98	69	301	0.78
2012-02-09	2012-02-09	19	337	351	0.78
2012-02-10	2012-02-10	52	193	652	0.78
2012-02-11	2012-02-11	133	63	199	0.78
2012-02-12	2012-02-12	87	216	310	0.78
2012-02-13	2012-02-13	4	81	199	0.78

2012-02-14	2012-02-14	0	115	469	0.78
2012-02-15	2012-02-15	112	348	278	0.78
2012-02-16	2012-02-16	155	220	69	0.78
2012-02-17	2012-02-17	176	20	560	0.78
2012-02-18	2012-02-18	50	24	314	0.78
2012-02-19	2012-02-19	122	73	369	0.78
2012-02-20	2012-02-20	111	8	595	0.78
2012-02-21	2012-02-21	24	68	259	0.78
2012-02-22	2012-02-22	46	195	398	0.78
2012-02-23	2012-02-23	64	180	19	0.78
2012-02-24	2012-02-24	115	334	343	0.78
2012-02-25	2012-02-25	4	282	337	0.78
2012-02-26	2012-02-26	115	64	396	0.78
2012-02-27	2012-02-27	48	264	347	0.78
2012-02-28	2012-02-28	58	5	670	0.78
2012-02-29	2012-02-29	38	176	445	0.78

	cost_of_banana	cost_of_cake	apples_profit	bananas_profit	\
date					
2012-02-01	0.48	1.07	66.25	19.88	
2012-02-02	0.49	1.07	138.75	136.29	
2012-02-03	0.49	1.07	28.07	65.00	
2012-02-04	0.49	1.07	148.92	39.30	
2012-02-05	0.49	1.07	39.77	85.40	
2012-02-06	0.49	1.07	53.81	128.59	
2012-02-07	0.49	1.07	100.62	74.74	
2012-02-08	0.49	1.07	76.45	33.49	
2012-02-09	0.49	1.07	14.82	163.59	
2012-02-10	0.49	1.08	40.57	93.70	
2012-02-11	0.49	1.08	103.79	30.59	
2012-02-12	0.49	1.08	67.90	104.89	
2012-02-13	0.49	1.08	3.12	39.34	
2012-02-14	0.49	1.08	0.00	55.86	
2012-02-15	0.49	1.08	87.45	169.06	
2012-02-16	0.49	1.08	121.03	106.89	
2012-02-17	0.49	1.08	137.45	9.72	
2012-02-18	0.49	1.08	39.05	11.66	
2012-02-19	0.49	1.08	95.30	35.48	
2012-02-20	0.49	1.08	86.72	3.89	
2012-02-21	0.49	1.08	18.75	33.06	
2012-02-22	0.49	1.08	35.95	94.81	
2012-02-23	0.49	1.08	50.02	87.53	
2012-02-24	0.49	1.08	89.89	162.44	
2012-02-25	0.49	1.08	3.13	137.17	
2012-02-26	0.49	1.08	89.91	31.13	
2012-02-27	0.49	1.08	37.53	128.44	
2012-02-28	0.49	1.08	45.36	2.43	

2012-02-29	0.49	1.08	29.72	85.65
------------	------	------	-------	-------

	cakes_profit	total_profit	Day	Year	Fiscal Year
date					
2012-02-01	291.01	377.14	Wednesday	2012	2012
2012-02-02	361.93	636.97	Thursday	2012	2012
2012-02-03	50.48	143.55	Friday	2012	2012
2012-02-04	239.56	427.77	Saturday	2012	2012
2012-02-05	337.35	462.52	Sunday	2012	2012
2012-02-06	204.16	386.56	Monday	2012	2012
2012-02-07	21.49	196.85	Tuesday	2012	2012
2012-02-08	323.51	433.45	Wednesday	2012	2012
2012-02-09	377.29	555.71	Thursday	2012	2012
2012-02-10	700.93	835.21	Friday	2012	2012
2012-02-11	213.96	348.34	Saturday	2012	2012
2012-02-12	333.35	506.14	Sunday	2012	2012
2012-02-13	214.01	256.48	Monday	2012	2012
2012-02-14	504.45	560.31	Tuesday	2012	2012
2012-02-15	299.05	555.55	Wednesday	2012	2012
2012-02-16	74.23	302.16	Thursday	2012	2012
2012-02-17	602.55	749.72	Friday	2012	2012
2012-02-18	337.90	388.62	Saturday	2012	2012
2012-02-19	397.14	527.92	Sunday	2012	2012
2012-02-20	640.45	731.06	Monday	2012	2012
2012-02-21	278.82	330.63	Tuesday	2012	2012
2012-02-22	428.51	559.27	Wednesday	2012	2012
2012-02-23	20.46	158.01	Thursday	2012	2012
2012-02-24	369.38	621.71	Friday	2012	2012
2012-02-25	362.97	503.26	Saturday	2012	2012
2012-02-26	426.57	547.61	Sunday	2012	2012
2012-02-27	373.83	539.81	Monday	2012	2012
2012-02-28	721.90	769.69	Tuesday	2012	2012
2012-02-29	479.53	594.90	Wednesday	2012	2012

We can also ‘slice’ a DataFrame with a DatetimeIndex using date strings. Let’s grab the data for the last four days of August 2014.

```
[85]: df.loc['2014-08-28':'2014-08-31'] # Last 4 days of August 2014
```

```
[85]:
```

	date	apples_sold	bananas_sold	cakes_sold	cost_of_apple \
date					
2014-08-28	2014-08-28	76	186	569	0.88
2014-08-29	2014-08-29	71	46	662	0.88
2014-08-30	2014-08-30	67	286	487	0.88
2014-08-31	2014-08-31	123	317	163	0.88

	cost_of_banana	cost_of_cake	apples_profit	bananas_profit \
--	----------------	--------------	---------------	------------------

date				
2014-08-28	0.55	1.21	66.58	101.38
2014-08-29	0.55	1.21	62.20	25.08
2014-08-30	0.55	1.21	58.71	155.93
2014-08-31	0.55	1.21	107.79	172.85

	cakes_profit	total_profit	Day	Year	Fiscal Year
date					
2014-08-28	686.75	854.71	Thursday	2014	2015
2014-08-29	799.10	886.38	Friday	2014	2015
2014-08-30	587.93	802.56	Saturday	2014	2015
2014-08-31	196.81	477.45	Sunday	2014	2015

To round things off, let's group the data by quarters of the fiscal year (which starts in April) and plot mean profit over time.

```
[111]: # A new column for quarters of the fiscal year
df['Fiscal Year'] = df['date'].dt.to_period('Q-APR')
df.head()
```

```
[111]:
```

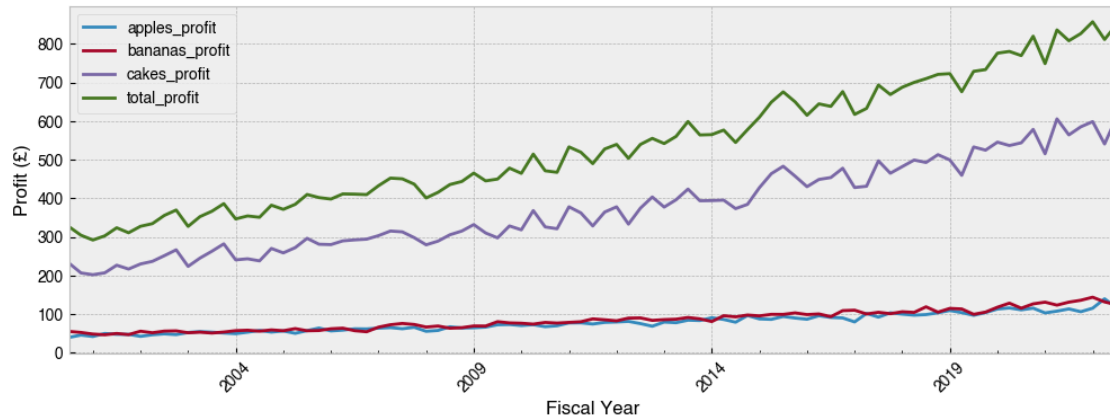
	date	apples_sold	bananas_sold	cakes_sold	cost_of_apple \
date					
2000-01-01	2000-01-01	175	174	534	0.45
2000-01-02	2000-01-02	61	316	394	0.45
2000-01-03	2000-01-03	164	320	487	0.45
2000-01-04	2000-01-04	55	291	143	0.45
2000-01-05	2000-01-05	8	261	630	0.45

	cost_of_banana	cost_of_cake	apples_profit	bananas_profit \
date				
2000-01-01	0.28	0.62	78.75	48.72
2000-01-02	0.28	0.62	27.45	88.49
2000-01-03	0.28	0.62	73.82	89.62
2000-01-04	0.28	0.62	24.76	81.51
2000-01-05	0.28	0.62	3.60	73.12

	cakes_profit	total_profit	Day	Year	Fiscal Year
date					
2000-01-01	331.08	458.55	Saturday	2000	2000Q3
2000-01-02	244.31	360.25	Sunday	2000	2000Q3
2000-01-03	302.02	465.46	Monday	2000	2000Q3
2000-01-04	88.69	194.96	Tuesday	2000	2000Q3
2000-01-05	390.79	467.51	Wednesday	2000	2000Q3

```
[112]: # Get the profit columns
profit_cols = df.columns[df.columns.str.endswith('profit')]
(
```

```
df.groupby('Fiscal Year')[profit_cols]
    .mean()
    .plot(ylabel='Profit (£)', figsize=(12, 4), rot=45)
);
```



If you can master tricks like this in **pandas**, the world of time series data is at your fingertips.

[Follow this link to learn more about Time Series / Date functionality in pandas](#)